

# Credit Card Fraud Detection


AM20216801  
장철희

# Index

- Examples of anomaly detection
- PreProcessing
- modeling
- Result
- Conclusion
- Reference

# Examples of anomaly detection

Cyber-intrusion




Fraud



Malware



Medical




Social Network




```
localhost CRON[19637]: (pam_unix) session closed for user root
localhost CRON[19013]: (pam_unix) session opened for user root by [uid=0]
localhost sashd[21417]: (pam_unix) authentication failure; logname= uid=0
localhost CRON[19013]: (pam_unix) session closed for user root
localhost CRON[20013]: (pam_unix) session opened for user root by [uid=0]
localhost sudo: user1: TTY=pts/2 : PWD=/home/user1 : USER=root : CON
localhost sashd[30504]: (pam_unix) authentication failure; logname= uid=0
localhost sashd[30504]: Accepted password for user1 from 10.15.1.39 port
localhost su[30694]: + pts/2 root:root
```

Log file



IoT Big-Data



Industrial



Video Surveillance

# Anomaly Detection

## **Supervised Anomaly Detection :**

**Supervised anomaly detection** techniques require a data set that has been labeled as "normal" and "abnormal" and involves training a classifier. However, this approach is rarely used in anomaly detection due to the general unavailability of labelled data and the inherent unbalanced nature of the classes.

## **Semi-supervised (One-Class) Anomaly Detection**

**Semi-supervised anomaly detection** techniques assume that some portion of the data is labelled. This may be any combination of the normal or anomalous data, but more often than not the techniques construct a model representing normal behavior from a given *normal* training data set, and then test the likelihood of a test instance to be generated by the model.

## **Unsupervised Anomaly Detection**

**Unsupervised anomaly detection** techniques assume the data is unlabelled and are by far the most commonly used due to their wider and relevant application.

# PreProcessing

[1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

#for data preprocessing
from sklearn.decomposition import PCA

#for modeling
from sklearn.neighbors import LocalOutlierFactor
from sklearn.ensemble import IsolationForest

#filter warnings
import warnings
warnings.filterwarnings("ignore")

import os
print(os.listdir("../input"))

# Any results you write to the current directory are saved as output.
```

['creditcard.csv']

# PreProcessing


```
[2]: df = pd.read_csv("../input/creditcard.csv")
df.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0

5 rows × 31 columns

```
[3]: sns.countplot(df.Class)
plt.show()
print(df.Class.value_counts())
```



```
0    284315
1         492
Name: Class, dtype: int64
```


Data that has been converted to V1 to V28 by PCA was used due to confidentiality issues.

# PreProcessing

[4]:

```
timedelta = pd.to_timedelta(df['Time'], unit='s')
df['Time_hour'] = (timedelta.dt.components.hours).astype(int)

plt.figure(figsize=(12,5))
sns.distplot(df[df['Class'] == 0]['Time_hour'], color='g')
sns.distplot(df[df['Class'] == 1]['Time_hour'], color='r')
plt.title('Fraud and Normal Transactions by Hours', fontsize=17)
plt.xlim([-1,25])
plt.show()
```



# PreProcessing

```
[5]: cols= df[['Time', 'Amount']]

pca = PCA()
pca.fit(cols)
X_PCA = pca.transform(cols)

df['V29']=X_PCA[:,0]
df['V30']=X_PCA[:,1]

df.drop(['Time','Time_hour', 'Amount'], axis=1, inplace=True)

df.columns

[5]: Index(['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11',
          'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21',
          'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'class', 'V29', 'V30'],
          dtype='object')
```

PCA.fit() : found the principal

PCA.transform() : Transform data into new principal components



# PreProcessing


[6]:

```
columns = df.drop('Class', axis=1).columns
grid = gridspec.GridSpec(6, 5)

plt.figure(figsize=(20,10*2))

for n, col in enumerate(df[columns]):
    ax = plt.subplot(grid[n])
    sns.distplot(df[df.Class==1][col], bins = 50, color='g')
    sns.distplot(df[df.Class==0][col], bins = 50, color='r')
    ax.set_ylabel('Density')
    ax.set_title(str(col))
    ax.set_xlabel('')

plt.show()
```



# PreProcessing

[7]:

```
def ztest(feature):  
  
    mean = normal[feature].mean()  
    std = fraud[feature].std()  
    zScore = (fraud[feature].mean() - mean) / (std/np.sqrt(sample_size))  
  
    return zScore
```

[8]:

```
columns= df.drop('Class', axis=1).columns  
normal= df[df.Class==0]  
fraud= df[df.Class==1]  
sample_size=len(fraud)  
significant_features=[]  
critical_value=2.58  
  
for i in columns:  
  
    z_vavalue=ztest(i)  
  
    if( abs(z_vavalue) >= critical_value):  
        print(i, " is statistically significant") #Reject Null hypothesis. i.e. H0  
        significant_features.append(i)
```

```
V1 is statistically significant  
V2 is statistically significant  
V3 is statistically significant  
V4 is statistically significant  
V5 is statistically significant  
V6 is statistically significant  
V7 is statistically significant  
V9 is statistically significant  
V10 is statistically significant  
V11 is statistically significant  
V12 is statistically significant  
V14 is statistically significant  
V16 is statistically significant  
V17 is statistically significant  
V18 is statistically significant  
V19 is statistically significant  
V20 is statistically significant  
V21 is statistically significant  
V24 is statistically significant  
V27 is statistically significant  
V28 is statistically significant  
V29 is statistically significant  
V30 is statistically significant
```

## z-test:

The z-test is an analytical technique that tests a hypothesis by comparing the means of two groups.

Valid transactions as our population

Fraud transactions as sample

Two tailed Z-test

Level of significance 0.01

Corresponding critical value is 2.58

Hypothesis:

H0: There is no difference (insignificant)

H1: There is a difference (significant)

$$Zscore = (\bar{x} - \mu) / S.E$$

# PreProcessing

```
[9]: significant_features.append('Class')
     df = df[significant_features]

     inliers = df[df.Class==0]
     ins = inliers.drop(['Class'], axis=1)

     outliers = df[df.Class==1]
     outs = outliers.drop(['Class'], axis=1)

     ins.shape, outs.shape
```

```
[9]: ((284315, 23), (492, 23))
```

```
[10]: def normal_accuracy(values):

       tp=list(values).count(1)
       total=values.shape[0]
       accuracy=np.round(tp/total,4)

       return accuracy

     def fraud_accuracy(values):

       tn=list(values).count(-1)
       total=values.shape[0]
       accuracy=np.round(tn/total,4)

       return accuracy
```

# Modeling


```
[11]: state= 42

ISF = IsolationForest(random_state=state)
ISF.fit(ins)

normal_isf = ISF.predict(ins)
fraud_isf = ISF.predict(outs)

in_accuracy_isf=normal_accuracy(normal_isf)
out_accuracy_isf=fraud_accuracy(fraud_isf)
print("Accuracy in Detecting Normal Cases:", in_accuracy_isf)
print("Accuracy in Detecting Fraud Cases:", out_accuracy_isf)

Accuracy in Detecting Normal Cases: 0.9
Accuracy in Detecting Fraud Cases: 0.9004
```



## IsolationForest

It is mainly used to detect outliers in the current data. As the name suggests, it is implemented based on a tree, and it is implemented by splitting the data at random and isolating all observations.

In particular, it has the advantage of being able to operate efficiently on data with many variables.

# Modeling

```
[12]: LOF = LocalOutlierFactor(novelty=True)
      LOF.fit(ins)

      normal_lof = LOF.predict(ins)
      fraud_lof = LOF.predict(outs)

      in_accuracy_lof=normal_accuracy(normal_lof)
      out_accuracy_lof=fraud_accuracy(fraud_lof)
      print("Accuracy in Detecting Normal Cases:", in_accuracy_lof)
      print("Accuracy in Detecting Fraud Cases:", out_accuracy_lof)
```

```
Accuracy in Detecting Normal Cases: 0.9171
Accuracy in Detecting Fraud Cases: 0.5142
```

## LocalOutlierFactor


LOF indicates how far each observation is within the data (outliers). The most important characteristic of LOF is not to consider all data as a whole, but rather to use the data surrounding that observation to identify the extent of an outlier from a local point of view.

# Result

```
[13]: fig, (ax1,ax2)= plt.subplots(1,2, figsize=[15,2])

ax1.set_title("Accuracy of Isolation Forest",fontsize=20)
sns.barplot(x=[in_accuracy_isf,out_accuracy_isf],
            y=['normal', 'fraud'],
            label="classifiers",
            color="b",
            ax=ax1)
ax1.set(xlim=(0,1))

ax2.set_title("Accuracy of Local Outlier Factor",fontsize=20)
sns.barplot(x=[in_accuracy_lof,out_accuracy_lof],
            y=['normal', 'fraud'],
            label="classifiers",
            color="r",
            ax=ax2)
ax2.set(xlim=(0,1))
plt.show()
```



# Conclusion

Both, Isolation Forest and Local Outlier Factor performed same in predicting Normal cases but Isolation Forest performed far better in detecting Fraud cases.

# Reference

<https://www.kaggle.com/code/sabanasimbutt/anomaly-detection-using-unsupervised-techniques>

“Deep Learning for Anomaly Detection: A Survey,” 2019 arXiv